

Lab 1 : Kubernetes cluster

Step 1: Prerequisites Setup

On node1

```
# yum install vim -y
# hostname master
# echo 127.0.0.1 master >> /etc/hosts
```

logout and login to check the hostname setup on the master machine

On node2

```
# yum install vim -y
# hostname worker1
# echo 127.0.0.1 worker1 >> /etc/hosts
```

logout and login to check the hostname setup on the worker1 machine

On node3

```
# yum install vim -y
# hostname worker2
# echo 127.0.0.1 worker2 >> /etc/hosts
```

logout and login to check the hostname setup on the worker2 machine

Step 2: Install Docker CE

Set up the repository and Install required packages.

```
# yum install yum-utils device-mapper-persistent-data lvm2
# yum-config-manager --add-repo https://download.docker.com/linux/centos/docker-ce.repo
```

Install Docker CE.

```
# yum update && yum install docker-ce-18.06.2.ce
```

Restart Docker

```
systemctl daemon-reload
systemctl restart docker
systemctl enable docker
```

Step 3: Install kubeadm

You will install these packages on all of your machines:

- **kubeadm**: the command to bootstrap the cluster.
- **kubelet**: the component that runs on all of the machines in your cluster and does things like starting pods and containers.
- **kubectl**: the command line util to talk to your cluster.

```
# cat <<EOF > /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=https://packages.cloud.google.com/yum/repos/kubernetes-el7-x86_64
enabled=1
gpgcheck=1
repo_gpgcheck=1
gpgkey=https://packages.cloud.google.com/yum/doc/yum-key.gpg
https://packages.cloud.google.com/yum/doc/rpm-package-key.gpg
EOF
```

```
# yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes
```

```
# systemctl enable --now kubelet
```

Step 4: Enable Net packet forwarding

Some users on RHEL/CentOS 7 have reported issues with traffic being routed incorrectly due to iptables being bypassed. You should ensure `net.bridge.bridge-nf-call-iptables` is set to 1 in your `sysctl` config, e.g.

```
# cat <<EOF > /etc/sysctl.d/k8s.conf
net.bridge.bridge-nf-call-ip6tables = 1
net.bridge.bridge-nf-call-iptables = 1
EOF
```

```
# sysctl -system
```

Step 5: Setup master node

The control-plane node is the machine where the control plane components run, including etcd (the cluster database) and the API server (which the kubectl CLI communicates with).

```
# kubeadm init --pod-network-cidr=192.16.0.0/16
```

To make kubectl work for your non-root user, run these commands, which are also part of the kubeadm init output:

```
# mkdir -p $HOME/.kube
# sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
# sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

Step 6: Enable Network CNI

We are using calico to setup the pod cluster network

```
# kubectl apply -f https://docs.projectcalico.org/v3.8/manifests/calico.yaml
```

Step 7: Adding worker node

Run the following command on the master node to generate command to add worker on the master node.

```
# kubeadm token create --print-join-command  
copy the output
```

Now login into each worker node and paste the command.

Once done, login back to the master node and run the following command to check the cluster status

```
# kubectl get nodes
```

You should be able to see all three nodes in ready state.

Congrats you have successfully deployed the kubernetes cluster using kubeadm.

Lab 2: Pods lab

Step 1: Namespace

Create a name space with the name new-ns with the manifest file

```
# vim ns.yml
```

```
apiVersion: v1
```

```
kind: Namespace
```

```
metadata:
```

```
  name: new-ns
```

save and quit the file

```
# kubectl create -f ns.yml
```

```
# kubectl get ns
```

Step 2: Pod

Create a pod name pod-data with the manifest file in the new-ns namespace with nginx image

```
# vim pod.yml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: pod1
```

```
  namespace: new-ns
```

```
spec:
```

```
  containers:
```

```
    - name: cont1
```

```
      image: nginx
```

save and quit the file

```
# kubectl create -f pod.yml
```

```
# kubectl get pods -n new-ns
```

Step 3: Pod with label

Create a pod named pod2 with label env=test with redis image

```
# vim pod.yml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: pod2
```

```
  labels:
```

```
    env: test
```

```
spec:
```

```
  containers:
```

```
    - name: cont1
```

```
      image: redis
```

save and quit the file

```
# kubectl create -f pod2.yml
```

```
# kubectl get pods -o wide
```

```
# kubectl get pods --show-labels
```

Lab 3: Replica's

Step 1: RC

Setup a replica controller rc1 with 3 replicas and the pod selector to be env=prod

```
# vim rc.yml
```

```
apiVersion: v1
```

```
kind: ReplicationController
```

```
metadata:
```

```
  name: rc1
```

```
spec:
```

```
  replicas: 3
```

```
  selector:
```

```
    env: prod
```

```
  template:
```

```
    metadata:
```

```
      labels:
```

```
        env: prod
```

```
    spec:
```

```
      containers:
```

```
      - name: cont1
```

```
        image: nginx
```

save and quit the file

```
# kubectl create -f rc.yml
```

```
# kubectl get rc
```

```
# kubectl get pods -o wide
```

Step 2: ReplicaSet

Create a Replica set rs1 with 3 replicas and two pod selection as env = test or env = prod

```
# vim rs.yml
```

```
apiVersion: apps/v1
```

```
kind: ReplicaSet
```

```
metadata:
```

```
  name: rs1
```

```
spec:
```

```
  replicas: 3
```

```
  selector:
```

```
    matchExpressions:
```

```
      - key: env
```

```
        operator: In
```

```
        values:
```

```
          - prod
```

```
          - test
```

```
  template:
```

```
    metadata:
```

```
      labels:
```

```
        env: prod
```

```
    spec:
```

```
      containers:
```

```
        - name: cont1
```

```
          image: nginx
```

save and quit the file

```
# kubectl create -f rs.yml
```

```
# kubectl get rs
```

```
# kubectl get pods
```

Step 3: DaemonSet

Create the daemonset and see it's operation

```
# vim ds.yaml
```

```
apiVersion: apps/v1
```

```
kind: DaemonSet
```

```
metadata:
```

```
  name: ds1
```

```
spec:
```

```
  selector:
```

```
    matchLabels:
```

```
      abc: xyz
```

```
  template:
```

```
    metadata:
```

```
      labels:
```

```
        abc: xyz
```

```
    spec:
```

```
      containers:
```

```
        - name: newcont
```

```
          image: nginx
```

save and quit the file

```
# kubectl create -f ds.yaml
```

```
# kubectl get daemonset
```

```
# kubectl get pods -o wide
```


Lab 3: Services

Step 1: ClusterIP Service

Create a pod with label new=old and nginx image and expose it with clusterIP service on port 8080

```
# vim l3pod.yml
```

```
apiVersion: v1
kind: Pod
metadata:
  name: l3pod
  labels:
    new: old
spec:
  containers:
    - name: cont1
      image: nginx
```

Save and quit the file

```
# kubectl create -f l3pod.yml
```

```
# kubectl get pods -o wide
```

```
# vim cip.yml
```

```
apiVersion: v1
kind: Service
metadata:
  name: cipservice
spec:
  type: ClusterIP
  selector:
    new: old
  ports:
    - name: port1
      port: 8080
      TargetPort: 80
```

save and quit the file

```
# kubectl create -f cip.yml
```

```
# kubectl get svc
```

```
# curl <serviceipaddress>
```

Step 1: NodePort Service

Create a pod with label old=new and nginx image and expose it with clusterIP service on port 8080 and node port 32080

```
# vim l3pody.yml
```

```
apiVersion: v1
kind: Pod
metadata:
  name: l3pody
  labels:
    old: new
spec:
  containers:
    - name: cont1
      image: nginx
```

Save and quit the file

```
# kubectl create -f l3pod.yml
```

```
# kubectl get pods -o wide
```

```
# vim nip.yml
```

```
apiVersion: v1
kind: Service
metadata:
  name: cipservice
spec:
  type: NodePort
  selector:
    new: old
  ports:
    - name: port1
      port: 8080
      TargetPort: 80
      nodePort: 32080
```

save and quit the file

```
# kubectl create -f nip.yml
```

```
# kubectl get svc
```

Open your desktop browser and access <http://mastereip:32080>

Lab 4: Storages

Step 1: Host path Volume

Create a host path volume to mount /data directory into /vishal of the container

```
# vim l4pod.yml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: pod3
```

```
spec:
```

```
  containers:
```

```
    - name: cont1
```

```
      image: nginx
```

```
      volumeMounts:
```

```
        - name: vol1
```

```
          mountPath: /vishal
```

```
  volumes:
```

```
    - name: vol1
```

```
      hostPath:
```

```
        path: /data
```

```
# kubectl create -f l4pod.yml
```

```
# kubectl get pods
```

Lab 5: Injecting Data

Step 1: Configmap

Create a config map with name qwe with key value pair to be v1=vishal

Export this variable in the pod l5pod with environment variable name as username

```
# kubectl create configmap qwe --from-literal=v1=vishal
```

```
# kubectl get configmap
```

```
# vim l5pod.yml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: l5pod
```

```
spec:
```

```
  containers:
```

```
    - name: cont1
```

```
      image: nginx
```

```
      env:
```

```
        - name: username
```

```
          valueFrom:
```

```
            configMapKeyRef:
```

```
              name: qwe
```

```
              key: v1
```

save and quit the file

```
# kubectl create -f l5pod.yml
```

```
# kubectl get pods -o wide
```

```
# kubectl exec -it pod l5pod
```

inside the pod # echo \$username

the output should be vishal

Step 2: Secret

Create a config map with name sec1 with key value pair to be pass=redhat

Export this variable in the pod l5pod with environment variable name as password

```
# kubectl create secret generic sec1 --from-literal=pass=redhat
```

```
# kubectl get secret
```

```
# vim l5pod1.yml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: l5pod1
```

```
spec:
```

```
  containers:
```

```
    - name: cont1
```

```
      image: nginx
```

```
      env:
```

```
        - name: username
```

```
          valueFrom:
```

```
            secretKeyRef:
```

```
              name: sec1
```

```
              key: pass
```

save and quit the file

```
# kubectl create -f l5pod.yml
```

```
# kubectl get pods -o wide
```

```
# kubectl exec -it pod l5pod
```

inside the pod # echo \$password

the output should be redhat

Lab 6: Deployment

Step 1: Deployment

Deploy a deployment dp1 with redis image and 3 replicas.

```
# vim l6dp.yml
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: dp1
spec:
  replicas: 3
  selector:
    matchLabels:
      app: redis
  template:
    metadata:
      labels:
        app: redis
    spec:
      containers:
        - name: cont1
          image: redis
```

save and quit the file

```
# kubectl create -f l6dp.yml
```

```
# kubectl get deployment
```

```
# kubectl get pods -o wide
```